

CCDetect-LSP: Incremental code clone detection for IDEs

Language- and IDE agnostic incremental code clone detection for IDEs

Jakob Konrad Hansen

60 ECTS study points

Department of Informatics
Faculty of Mathematics and Natural Sciences

Jakob Konrad Hansen

CCDetect-LSP: Incremental code clone detection for IDEs

Language- and IDE agnostic incremental code
clone detection for IDEs

Abstract

TODO: Write abstract

Contents

1	Introduction	4
1.1	Motivation and problem statement	4
1.2	Structure	5
1.3	Our contribution	5
2	Background	6
2.1	Software quality and duplicated code	6
2.2	Code clones	6
2.3	Code clone detection process and techniques	8
2.4	Incremental clone detection	10
2.5	Code clone IDE tooling	10
2.6	The Language Server Protocol	11
2.7	Preliminary algorithms and data structures	12
3	Implementation: LSP server architecture	19
3.1	Document index	20
3.2	Displaying and interacting with clones	21
4	Implementation: Initial detection	24
4.1	Fragment selection	24
4.2	Fingerprinting	25
4.3	Suffix array construction	26
4.4	Clone extraction	30
4.5	Source-mapping	31
5	Implementation: Incremental detection	35
5.1	Affordable operations	35
5.2	Updating the document index	35
5.3	Updating fingerprints	36
5.4	Computing edit operations	37
5.5	Dynamic suffix arrays	43
5.6	Storing old clones	43
6	Evaluation	45
6.1	Big-O analysis of incremental detection	45
6.2	Performance comparison	45

6.3	Comparison with iClones	45
6.4	Verifying clones with BigCloneBench	45
7	Discussion	46
7.1	Speed of incremental updates vs linear-time SACA	46
7.2	Comparison with iClones	46
7.3	Chosen clones	46
7.4	Usability while programming	46
8	Conclusion and future work	47
8.1	Future work	47
8.2	Conclusion	47
9	Appendix	48

Chapter 1

Introduction

Refactoring is the process of restructuring source code in order to improve the internal behavior of the code, without changing the external behavior [10, p. 9]. Refactoring source code is often performed in order to eliminate instances of bad design quality in code, otherwise known as code smells.

A study conducted by Diego Cedrim et al. has shown that while programmers tend to refactor smelly code, they are rarely successful at eliminating the smells they are targeting [8]. Most refactorings performed were either “smell-neutral”, meaning that the code smell is not eliminated, or “stinky”, meaning that they introduced more code smells than they eliminated. Automated tools that help programmers make better refactorings and perform code analysis, could be a solution to this problem.

Duplicated code is a code smell which occurs in practically every large software project. Code clone analysis (duplicate code analysis) has in the last decade become a highly active field of research and many tools have been developed to detect duplicated code [15, p. 6].

1.1 Motivation and problem statement

Many tools and algorithms exist for duplicate code detection. However, few of these have the capability of efficiently detecting duplicated code in a real-time IDE environment. Incremental algorithms that do not recompute all clones from scratch are interesting for use-cases such as IDEs and different Git revisions of the same source code, but this type of algorithm has not been thoroughly explored for code clone analysis.

The current landscape of clone detection algorithms and tools is therefore lacking in terms of integration and support for fast incremental updates within an IDE. This limits the ability for programmers to easily detect duplicated code as they work, as these tools are also often limited to specific programming languages or IDEs.

Our proposed solution addresses this issue by introducing a new algorithm that is capable of detecting and updating all existing code clones in real-time as source code changes, faster than redoing the analysis from scratch. The tool which implements this algorithm is also programming

language and IDE agnostic, which allows programmers to seamlessly incorporate the detection of duplicated code into their existing development environment.

1.2 Structure

The current chapter describes the motivation and contribution of the thesis. Chapter 2 continues by giving some background on code clone analysis, some existing tools and some preliminary algorithms and data structures used in the implementation. Chapter 3-5 describes the implementation of the tool and the algorithms used for initial and incremental clone detection. In chapter 6, the tool is evaluated based on different criteria, and compared to other existing solutions. Chapter 7 discusses the results of the evaluation and discusses some choices made in the implementation. Chapter 8 concludes the thesis and lists future work.

1.3 Our contribution

This thesis will present and evaluate a tool which provides clone detection capabilities in a real-time IDE environment. The main goal is to create a tool which fits well into the development cycle and can efficiently update its analysis while writing code.

The tool will allow the user to list and interact with all the code clones in the codebase, jump between matching clones, and get quick feedback while editing code in order to determine which clones are introduced/eliminated.

Existing incremental clone detection tools either do not fit into an IDE scenario (techniques based on distributed computing), are limited in terms of what clones they display, or have not been shown to scale well in terms of processing time or memory usage for larger codebases. Therefore, this thesis will focus on the following areas.

Incremental clone detection: The main focus of this thesis is making the tool efficient in terms of incrementally updating whenever edits are performed in the IDE. Most clone detection tools calculate clones from scratch and have no functionality to more efficiently calculate the clones after a small edit has been applied to the source code. Our tool implements a novel application of dynamic suffix arrays to quickly update and add/remove clones, often faster than calculating the clones from scratch using a linear time suffix array construction algorithm.

IDE and language agnostic clone detection: The tool gives programmers the ability to view clones in their IDE. Utilizing features of the language server protocol (LSP) such as diagnostics and code-actions, the tool provides clone analysis to any editor which implements the LSP protocol. As far as we are aware there are no other clone detection tools which utilize the LSP protocol to provide clone analysis. The tool is also language agnostic in the sense that it only needs a grammar for a language to analyze it.

Chapter 2

Background

2.1 Software quality and duplicated code

Software quality is hard to define. The term “quality” is ambiguous and is in the case of software quality, multidimensional. Quality in itself has been defined as “conformance to requirements” [9, p. 8]. In software, a simple measure of “conformance to requirements” is correctness, and a lack of bugs. However, software quality is often measured in other metrics, including metrics which are not directly related to functionality [18, p. 29]. These metrics often include maintainability, analyzability and changeability.

These metrics are affected negatively by duplicated code, code which is more or less copied to different locations in the source code. Multiple studies have shown that software projects typically contains 10 – 15% duplicated code [2]. Therefore, research into tools and techniques which can assist in reducing duplicated code will be of benefit to almost all software.

Duplicated code can lead to a plethora of antipatterns, and will often lead to an increase in technical debt. Technical debt occurs when developers make technical compromises that are expedient in the short term, but increases complexity in the long-term [3, p. 111]. An example of this, in the context of duplicated code, is the “Shotgun-Surgery” [10, p. 66] antipattern. This antipattern occurs when a developer wants to implement a change, but needs to change code at multiple locations for the change to take effect. This is a typical situation which slows down development and reduces maintainability when the amount of duplicated code increases in a software project.

2.2 Code clones

Duplicated code is often described as “code clones”, as a pair of code snippets which are duplicated are considered clones of each other.

Definition 1 (Code snippet). *A code snippet is a piece of contiguous source code in a larger software system.*

Definition 2 (Code clone). *A code clone is a code snippet which is equal to or similar to another*

code snippet. The two code snippets are both code clones, and together they form a code clone pair. Similarity is determined by some metric such as number of equal lines of code.

Definition 3 (Clone class). *A clone class is a set of code snippets where all snippets are considered clones of each other.*

The clone relation is a relation between code snippets which defines pairs of clones. The clone relation is reflexive and symmetric, but not necessarily transitive. The transitive property depends on the threshold for similarity when identifying code clones. Given

$$a \xleftrightarrow{\text{clone}} b \xleftrightarrow{\text{clone}} c$$

where a, b, c are code snippets and $\xleftrightarrow{\text{clone}}$ gives the clone relation, a is a clone of b , but not necessarily similar enough to be a clone of c , depending on the threshold for similarity. If the threshold for similarity is defined such that only equal clones are considered clones, the relation becomes transitive, and equivalence classes form clone classes.

Code clones are generally classified into four types [15]. The types classify code snippets as code clones with an increasing amount of leniency. Therefore, Type-1 code clones are very similar, while Type-4 clones are not necessarily syntactically similar at all. When defining types, it is the syntactic and structural differences which is compared, not functionality. The set of code clones classified by a code clone type is also a subset of the next type, meaning all Type-1 clones are also Type-2 clones, but not vice versa.

The code clone types are defined as follows:

Type-1 clones are syntactically identical. The only differences allowed are elements without meaning, like comments and white-space. Example:

TODO: Redo these

<pre> for (int i = 0; i < 10; i++) { print(i); } </pre>	<pre> for (int i = 0; i < 10; i++) { // A comment print(i); } </pre>
--	---

Type-2 clones are structurally identical. Possible differences include identifiers, literals and types. Type-2 clones are not much harder to detect than type-1 clones, since consistently renaming identifiers, literals and types allow a type-1 detection algorithm to find type-2 clones [35]. This type of clone is relevant to consider in refactoring scenarios when merging code clones. This is because type-2 clones can be relatively simple to parameterize in order to hide the differences in the merged code, and allows the original locations of clones to be replaced with a call to the merged code, with different parameters. Example:

<pre>for (int i = 0; i < 10; i++) { print(i); }</pre>	<pre>for (int (*\textbf j*) = (*\textbf 1*); (*\textbf j*) < 10; (*\textbf j*)++) { print((*\textbf j*)); }</pre>
--	--

Type-3 clones are required to be structurally similar, but not equal. Differences include statements that are added, removed or modified. This clone type relies on a threshold θ which determines how structurally different snippets can be in order to be considered Type-3 clones [15]. The granularity for this difference could be based on differing tokens, lines, etc. Detecting this type of clone is hard. Example:

<pre>// Clone 1 for (int i = 0; i < 10; i++) { print(i); }</pre>	<pre>// Clone 2 for (int i = 0; i < 10; i++) { print(i); (*\textbf{\int x = 10;}\textbf{*}) }</pre>
---	--

In this example there is a one line difference between the two snippets, so if $\theta \geq 1$, the two snippets would be considered Type-3 clones.

Type-4 clones have no requirement for syntactical or structural similarity, but are generally only relevant to detect when they have similar functionality. Detecting this type of clone is very challenging, but attempts have been made using program dependency graphs [34]. The following example shows two code snippets which have no clear syntactic or structural similarity, but is functionally equal:

<pre>print((n*(n-1))/2)</pre>	<pre>int sum = 0; for (int i = 0; i < n; i++) { for (int j = i+1; j < n; j++) { sum++; } } print(sum);</pre>
-------------------------------	--

Type-1 clones are often referred to as “exact” clones, while Type-2 and Type-3 clones are referred to as “near-miss” clones [35, p. 1].

2.3 Code clone detection process and techniques

The **Code clone detection process** is generally split into (but is not limited to) a sequence of steps to identify clones [15]. This process is often a pipeline of input-processing steps before finally comparing fragments against each other and filtering. The steps are generally as follows:

1. **Pre-processing:** Filter uninteresting code that we do not want to check for clones, for example generated code. Then partition code into a set of fragments, depending on granularity such as methods, files or lines.
2. **Transformation:** Transform fragments into an intermediate representation, with a source-

map back to the original code. An algorithm could potentially do multiple transformation before arriving at the wanted representation

- (a) **Extraction:** Transform source code into the input for the comparison algorithm. Can be tokens, AST, dependency graphs, suffix tree, etc.
 - (b) **Normalization:** Optional step which removes superficial differences such as comments, whitespace and identifier names. Often useful for identifying type-2 clones.
3. **Match detection:** Perform comparisons which outputs a set of candidate clone pairs.
 4. **Source-mapping / Formatting:** Convert candidate clone pairs from the transformed code back to clone pairs in the original source code.
 5. **Post-processing / Filtering:** Ranking and filtering manually or with automated heuristics
 6. **Aggregation:** Optionally aggregating sets of clone pairs into clone sets

Matching techniques are techniques which can be applied to source-code to detect clone-pairs. Most matching technique will also require specific pre-processing to be done in the earlier steps, for example building an AST. Some of the most explored techniques are as follows [26]:

Text-based approaches do little processing of the source code before comparing. Simple techniques such as fingerprinting or incremental hashing have been used in this approach. Dot-plots have also been used in newer text-based approaches, placing the hashes of fragments in a dot plot for use in comparisons.

Token-based approaches transform source code into a stream of tokens, similar to lexical scanning in compilers. The token stream is then scanned for duplicated subsequences of tokens. Since token streams can easily filter out superficial differences such as whitespace, indentation and comments, this approach is more robust to such differences. Concrete names of identifiers and values can be abstracted away when comparing the token-stream, therefore Type-2 clones can easily be identified. Type-3 clones can also be identified by comparing the fragments tokens and keeping clone pairs with a lexical difference lower than a given threshold. This can be solved with dynamic programming [4]. A common approach to detect clones using token-streams is with a suffix-tree [35]. A suffix-tree can solve the *Find all maximal repeats* problem efficiently, which in essence is the same problem as finding clone pairs. A similar algorithm can also be performed using a suffix-array instead, which requires less memory. This type of code clone detection is very fast compared to more intricate types of matching techniques.

Syntactic approaches transform source code into either concrete syntax trees or abstract syntax trees and find clones using either tree matching algorithms or structural metrics. For tree matching, the common approach is to find similar subtrees, which are then deemed as clone pairs. One way of finding similar subtrees is to compare subtrees with a tolerant tree matching algorithm for detecting type-3 clones [5]. Variable names, literal values and types may be abstracted to find type-2 clones more easily. Metrics-based techniques gather metrics for code fragments in the tree and uses the metrics to determine if the fragments are clones or not. One way is to use fingerprint functions where the fingerprint includes certain metrics, and compare the fingerprints of all fragments to find clones [17].

Hybrid approaches combine multiple approaches in order to improve detection. For example Zibran et al. developed a hybrid algorithm combining both token-based suffix trees for Type-1 and Type-2 clone detection, with a k-difference dynamic programming algorithm for Type-3 clone detection [35].

2.4 Incremental clone detection

Incremental clone detection involves avoiding recalculation of already calculated results when performing code clone detection. Since most code of a codebase will not change between revisions, a lot of processing can be avoided. However, this is not a simple problem, since changes in a single location can affect clone detection results across the entire codebase.

In order to incrementally detect code clones, an algorithm is run which calculates the initial clones, and for successive revisions of the source code, this list is incrementally updated, more efficiently than the initial run. Different algorithms will also maintain different data structures to support detecting new clones faster in successive revisions.

Göde and Koschke proposed the first incremental clone detection algorithm [12]. The algorithm employs a generalized suffix tree in which the amount of work of updating is only dependent on the size of the edited code. This approach requires a substantial amount of memory, and is therefore limited in scalability.

Nguyen et al. [22] showed that an AST-based approach utilizing “Locality-Sensitive Hashing” can detect clones incrementally with high precision, and showed that incremental updates could be done in real-time (< 1 second) for source code with a size of 300 KLOC.

Hummel et al. [14] later introduced the first incremental, scalable and distributed clone detection technique for Type-1 and Type-2 clones. This approach utilizes a custom “clone index” data structure which can be updated efficiently. The implementation of this data structure is similar to that of an inverted index. This technique uses distributed computing to speed up its detection process.

More recently, Ragkhitwetsagul and Krinke [24] presented the tool “Siamese”, which uses a novel approach of having multiple intermediate representations of source code to detect a high number of clones with support for incremental detection. The tool can detect up to Type-3 clones, but will only return clones based on “queries” given to it by the user. Queries are either files or methods in source-code, which are then checked for existing code clone.

2.5 Code clone IDE tooling

Developers are not always aware of the creation of clones in their code. *Clone aware development* means including clone management as a part of the software development process. Clone aware development therefore requires programmers to be aware of and be able to identify code clones while programming. Since large software projects can contain a lot of duplication, it can be hard to keep track of and manage clones. Tools which help developers locate and deal with clones can be a solution to this. However, Mathias Rieger et al. claims that a problem with many detection tools is that the tools “report large amounts data that must be treated with little tool support” [25, p. 1]. Detecting and eliminating clones early in their lifecycle with IDE integrated tools could be a solution to the problem of dealing with too many clones.

There are multiple existing clone detection tools, and the following section will go over tools that are integrated into an IDE and offer services to the programmer while developing. Some of these tools will

The IDE-based tools which exist can be categorized as follows [29, p. 8]:

- *Copy-paste-clones*: This category of tools deals only with code snippets which are copy-pasted from another location in code. These tools therefore only track clones which are created when copy-pasting, and does not use any other detection techniques. Therefore, this type of tool is not suitable for detecting clones which are made accidentally, since developers are aware that they are creating clones when pasting already existing code snippets.
- *Clone detection and visualization tools*: This category of tools has more sophisticated clone detection capabilities and will detect code clones which occur accidentally.
- *Versatile clone management*: This category of tools covers tools which provide more services than the above. Services like refactoring and simultaneous editing of clones fall under this category.

The following tools have been developed as IDE tools to allow for clone-aware development:

- Minhaz et al. introduced a hybrid technique for performing real-time focused searches, i.e. searching for code clones of a selected code snippet. This technique can also detect Type-3 clones [35]. It was later used in the tool *SimEclipse* [29] which is a plugin for the Eclipse editor. Since this tool can only detect clones of a code snippet which the developer actively selects, this tool is not well suited for finding accidental clones and tracking clones in areas.
- Another tool, SHINOBI, which is a plugin for the Visual Studio editor, can detect code clones in real-time without the need of the developer to select a code snippet, but the clones being displayed seem to only be clones of the source-code which is currently being edited. It can detect Type-1 and Type-2 code clones and uses a token-based suffix array approach to detect clones. The paper does not describe how the suffix array is updated, but as the paper was released before the first paper describing dynamic suffix arrays [27], one can assume that the suffix array is not dynamically updated. [20].
- The modern IDE IntelliJ has a built-in duplication detection and refactoring service, it is able to detect Type-1 and (some) type-2 code clones at a method granularity and refactors by replacing one of the clones with a method call to the other.

2.6 The Language Server Protocol

Static analysis tools that integrate with IDEs are usually tightly coupled to a specific IDE and its APIs, like parsing and refactoring support. This makes it difficult to utilize a tool in another IDE, since the API's the tool utilizes is no longer available. In order to make IDE-based static analysis tools more widely available, it is useful if such tools could be made IDE agnostic. The Language Server Protocol (LSP) is a protocol which attempts to solve this problem [21].

LSP is a protocol which specifies interaction between a client (IDE) and server in order to provide language tooling for the client. The goal of the protocol is to avoid multiple implementations of



Figure 2.1: Example client-server interaction using LSP

the same language tools for every IDE and every language, allowing for editor agnostic tooling. Servers which implement LSP will be able to offer IDEs code-completion, error-messages, go-to-definition and more. LSP also specifies generic code-actions and commands, which the LSP server provides to the client in order to perform custom actions defined by the server.

Figure 2.1 shows a sample interaction between client and server using LSP. The client sends requests to a server in the form of JSON-RPC messages, and the server sends a corresponding response, also in the form of JSON-RPC messages.

2.7 Preliminary algorithms and data structures

The following algorithms and data structures will be useful insight in following chapters to define the detection algorithm

Incremental-parsing

While writing code, programmers usually only edit small regions of code at a time. One “edit” will therefore only affect small parts of the internal representations of the code which most tools use to perform analysis. Reusing parts of this representation would therefore be faster and allow programming tools to scale better.

In our detection algorithm we will need to be able to parse our source code into an AST representation in order to select specific syntactic regions and extract the tokens. Therefore, we need a parsing algorithm, and when source code is edited we need an incremental parsing algorithm to more efficiently parse the source code when only a small region has been changed.

Incremental parsing is the process of reparsing only parts of a syntax tree whenever an edit is

performed. The motivation behind incremental parsing is to have a readily available syntax tree after every edit, while doing as little computing as possible to maintain it.

Tree-sitter is a parser generator tool which specializes in incremental parsing. It supports incremental parsing, error recovery and querying for specific nodes and subtrees [6]. These features combined allow Tree-sitter to become a powerful tool for editing and has been used for IDE features such as syntax-highlighting, refactoring and code navigation.

The incremental reparsing in Tree-sitter is performed using Wagner’s algorithm [33].

TODO: Explain Wagner’s algorithm with figure and example

Suffix trees

A classic algorithm [35, 12] for code clone detection traverses a suffix-tree in order to find maximal repeats in all suffixes of the input string T .

The suffix tree of a string T is a compressed trie where all the suffixes of T have been inserted. The tree is compressed by combining consecutive nodes in a row which has only one child into a single node. A common usage of a suffix tree is to determine whether or not a suffix exists in T , and where in T the suffix starts.

Figure 2.2 shows the suffix-tree for $T = \text{BANANA\$}$. In order to determine where the suffix $\text{ANANA\$}$ exists in T , one can start from the root, and traverse the tree, choosing the child node which correspond to the next character of the suffix which has not been “matched” yet.

$$\text{root} \xrightarrow{A} \text{node} \xrightarrow{NA} \text{node} \xrightarrow{NA\$} 1$$

Following this path, we see that the suffix $\text{ANANA\$}$ exists in T at index 1.

Suffix trees can be constructed in linear time with Ukkonen’s algorithm which builds a larger and larger suffix tree by inserting characters one by one and utilizing some tricks to avoid inserting suffixes before it needs to, lowering the complexity [31].

This data structure also facilitates solving the maximal repeat problem. A repeat in a string T is a substring that occurs at least twice in T . A maximal repeat in T is a repeat which is not a substring of another repeat in T , meaning that the maximal repeat cannot be extended in any direction to form a bigger repeat. The problem of finding all maximal repeats can be solved with a suffix tree using the following theorem:

Theorem 1 (Repeats in suffix tree). *Every internal node (except for the root) in a suffix tree corresponds to a substring which is repeated at least twice in T . The substring is found by concatenating the strings found on the path from the root of tree to the internal node.*

This theorem is explained by the fact that any internal node has at least two children, and a node having two children means that two suffixes share the same prefix up to that point. An algorithm which finds the maximal repeats would find the internal nodes which represents the longest strings.

The classic algorithm [35, 12] in terms of finding duplication in a string (such as source code)



Figure 2.2: Suffix tree for $T = \text{BANANA}\$$

using suffix trees would find all repeats of length k , where k is the threshold for how long a clone needs to be. This can be found by traversing the suffix tree and looking at all internal nodes which represent a string of length $\geq k$. Every internal node which represents a string which is $\geq k$ would correspond to a substring of the source code which occurs at least twice. Finding where the duplication occurs can be done by finding all the leaves of the subtree rooted in the internal node, which each hold the position where the suffix starts in T . Since a substring can have multiple repeats of different lengths longer than k , different strategies can be used to select which substrings are selected or not, such as filtering out repeats which are not maximal or repeats which contain or overlap each other.

In figure 2.2, there are three repeats, ANA, A and NA. The only maximal repeat would be ANA, since A and NA are not maximal.

Suffix arrays

The suffix array (SA) of a string T contains a lexicographical sorting of all suffixes in T . The suffix array does not contain the actual suffixes, but it contains integers pointing to the index where the suffix starts in T . Conversely, the inverse suffix array (ISA) contains integers describing which rank the suffix at a given position has. The rank of a suffix is its lexicographical ordering in T . ISA is therefore the inverse array of SA, such that if $\text{SA}[i] = n$, then $\text{ISA}[n] = i$.

Definition 4 (Suffix array). *Let T be a text of length N . The suffix array SA of T is an array of length N where $\text{SA}[i] = n$ if the suffix at $T[n..N-1]$ is the i th lexicographically smallest suffix in T .*

Definition 5 (Inverse suffix array). *Let T be a text of length N . The inverse suffix array ISA of T is an array of length N where $\text{ISA}[i] = n$ if the suffix at $T[i..N-1]$ is the n th smallest suffix in T lexicographically.*

The Longest-common prefix (LCP) array describes how many common characters there are in a prefix between two suffixes which are next to each other in the suffix array. Since the suffix array represents suffixes in a sorted order, the prefix length between adjacent suffixes in SA will be the longest possible common prefix for each suffix. These are the values in the LCP array.

Definition 6 (LCP array). *Let T be a text of length N and SA be the suffix array of T . The LCP array of T is an array of length N where $\text{LCP}[i] = n$ if the suffix $T[\text{SA}[i]..N]$ and $T[\text{SA}[i-1]..N]$*

Index	Suffix	Index	Suffix	Index	SA	ISA	LCP
0	BANANA\$	6	\$	0	6	4	0
1	ANANA\$	5	A\$	1	5	3	0
2	NANA\$	3	ANA\$	2	3	6	1
3	ANA\$	1	ANANA\$	3	1	2	3
4	NA\$	0	BANANA\$	4	0	5	0
5	A\$	4	NA\$	5	4	1	0
6	\$	2	NANA\$	6	2	0	2

(a) Suffixes (b) Sorted suffixes (c) SA, ISA and LCP

Table 2.1: $T = \text{BANANA\$}$

has a maximal common prefix of length n . $LCP[0]$ is undefined or 0.

For example, in table 2.1, the LCP value at position 3 contains the number of characters in the longest-common prefix of suffixes starting at position 1 (ANANA\$) and position 3 (ANA\$). These suffixes have 3 common characters in their prefix, therefore the LCP value at position 3 is 3.

Suffix arrays can be constructed in linear time in terms of the length of T . Many suffix array construction algorithms (SACA) have been discovered in the last decade [19], many of which run in linear-time. An algorithm which has been shown to be very efficient in practice is Nong and Chan’s algorithm based on recursive suffix sorting of smaller strings [23].

The enhanced (extended) suffix array is the suffix array and the additional LCP array, which has been shown to be as powerful as a suffix tree, in terms of what can be computed with it with the same time complexity, but uses a smaller amount of memory [1].

Dynamic bitsets

A bitset is an array of bits, each bit representing either the value true or false. A bit with the value of 1 is usually referred to as a set bit, and a value of 0 is referred to as an unset or cleared bit. Bitsets have at least operations for setting the value at a position, and looking up the value at a position. Bitsets are useful for many problems, especially as a “succinct data structure”. A succinct data structure is a data structure which attempts to use an amount of memory close to the theoretic lower bound, while still allowing effective queries on it. For example for a string of length n , we could store up to $O(n \log \sigma)$ bits before the bit vector exceeds the size of the string itself, where σ is the size of the string’s alphabet.

The most well known query to do on bitsets is the rank/select queries.

Definition 7 (Rank query). *A rank query $\text{rank}_1(i)$ on a bitset B , returns the number of set bits up to, but not including position i . Conversely, $\text{rank}_0(B)$ returns the number of unset bits up to i .*

Definition 8 (Select query). *A select query $\text{select}_1(i)$ on a bitset B , returns the position of the i th set bit in B . Conversely, $\text{select}_0(i)$ returns the position of the i th unset bit in B .*

Jacobson’s rank can calculate rank and select on static bitsets in $O(1)$ time by pre-calculating



Figure 2.3: Dynamic bitset

all answers in a space efficient table [16].

Definition 9 (Dynamic bitset). *A dynamic bitset is a bitset which in addition to other operations allow inserting and deleting bits (indel operations).*

An insert operation $\text{insert}(i, v)$ on a bitset B inserts the value v at position i in B , pushing all values at position $\geq i$ one position up.

A delete operation $\text{delete}(i)$ on a bitset B removes the value at position i in B , pushing all values at position $> i$ one position down.

The standard implementation of a dynamic bitset would implement the whole bitset as a single array of bytes B , which allows for accessing values in $O(1)$ time, but inserting and deleting takes $O(n)$ time, where n is the number of bits in B .

One way to speed up the indel operations would be to represent the bitset as a balanced tree containing multiple smaller bitsets. To represent a bitset of n bits, we can divide the bits into smaller bitsets, such that $O(\frac{n}{\log(n)})$ bitsets each contains $O(\log(n))$ bits. Since the bitsets are now of size $O(\log(n))$, inserting and deleting takes only $O(\log(n))$ time. The bitsets reside at the leaves of our balanced tree, and internal nodes contain only two integers, storing the number of bits in the left subtree (N), and the number of set bits in the left subtree (S). To access, insert or delete on index i , the tree is traversed to find the correct bitset where the i th bit is located, where the operation is done at that position. Finding the correct bitset and position is done by utilizing N and S in each node which is traversed. All operations now run in $O(\log(n))$ time, since traversing the tree to the correct bitset takes $O(\log(\frac{n}{\log(n)}))$ and performing the operation takes $O(\log(n))$ time.

Figure 2.3 shows how a dynamic bitset tree is structured, and Algorithm 1 shows how to access a value in the tree. Traversing the tree to calculate rank and select queries can be done similarly by summing set bits in left-subtrees (rank) or selecting which subtree to descend based on the number of set bits (select).

```

Algorithm access(node, i)
  if isLeaf(node) then
    | return node.bitset[i]
  end
  if node.N ≤ i then
    | return access(node.left, i)
  end
  return access(node.right, i - node.N)

```

Algorithm 1: Accessing a value in a dynamic bitset

Index	Cyclic-shift	Index	Cyclic-shift	L	F
0	BANANA\$	6	\$BANANA	0	$Rank_A(0) + C[A] = 0 + 1 = 1$
1	ANANA\$B	5	ABANAN	1	$Rank_N(1) + C[N] = 0 + 5 = 5$
2	NANA\$BA	3	ANAN\$B	2	$Rank_N(2) + C[N] = 1 + 5 = 6$
3	ANAN\$BAN	1	ANANA\$B	3	$Rank_B(3) + C[B] = 0 + 4 = 4$
4	NA\$BANA	0	BANANA\$	4	$Rank_{\$}(4) + C[\$] = 0 + 0 = 0$
5	A\$BANAN	4	NA\$BANA	5	$Rank_A(5) + C[A] = 1 + 1 = 2$
6	\$BANANA	2	NANA\$BA	6	$Rank_A(6) + C[A] = 2 + 1 = 3$
(a) Cyclic shifts		(b) Sorted cyclic shifts and BWT		(c) LF-mapping	

Table 2.2: $S = \text{BANANA\$}$, $\text{BWT} = \text{ANNB\$AA}$

Burrows-Wheeler transform

The Burrows-Wheeler transform (BWT) is a transform on strings, often performed on strings to improve compression [7]. The transform is computed by sorting all “cyclic-shifts” of the string lexicographically and extracting a new string from the last column of the cyclic-shift matrix. \$ is added to the string as a unique terminating character, this makes some of the algorithms on the BWT simpler. \$ is always the smallest character lexicographically. Table 2.2 shows the BWT for the string $S = \text{BANANA\$}$.

There is a direct correlation between the suffix array of S and the BWT of S . Table 2.2 also shows that the indices of the sorted cyclic-shifts correspond to exactly the SA of S . This coincides because when sorting cyclic-shifts, everything that occurs after the \$ of the cyclic-shift will not affect its lexicographical ordering. This is because no cyclic-shift will have a \$ in the same position, so comparing two cyclic shifts lexicographically will always terminate whenever a \$ is found. This means that the ordering of cyclic shifts is essentially the same as sorting all suffixes of S . Algorithm 2 shows how the BWT of S can be calculated directly from the SA of S in linear time. Since this is a one to one correlation between the SA and BWT, dynamic updates to a BWT would correspond to similar dynamic updates in the SA (and ISA). This will be useful in the following implementation chapter for the dynamic code clone detection algorithm.

An essential property of the BWT is that the transformation is reversible. By examining the BWT of a string T , we see that the BWT is a permutation of T . We will also see that there is a correlation between the characters in the first column of the cyclic-shift matrix and the last column (the BWT).

Algorithm ComputeBWT(S, SA)

```

     $n \leftarrow S.len$ 
     $BWT \leftarrow$  string of length  $n$ 
    for  $i$  from 0 to  $n$  do
        |  $pos \leftarrow (SA[i] - 1) \% n$ 
        |  $BWT[i] = S[pos]$ 
    end
    return BWT

```

Algorithm 2: Calculating the BWT of a string T from its suffix array

T can be calculated from the BWT as long as there is a unique terminating character ($\$$), or the position of the final character is stored. T is computed backwards by starting at the final character ($T[n - 1]$) and then finding the cyclic-shift where that character occurs in the first column (the previous cyclic-shift). The character in the last column of that cyclic-shift is $T[n - 2]$. If there are multiple of the same character, the n th occurrence of a certain symbol in the last column will map to the n th occurrence of the same symbol in the first column. This process is repeated until we finish the cycle, returning to the final character in the last column. Essentially, this process consists of traversing cyclic-shifts backwards and looking at the final character, which will give us T , since the final character of the cyclic-shift is continually shifting one position.

We can use the fact that the first column consists of the letters of T in sorted order to determine the previous cyclic-shift with only the last column. We can calculate the previous cyclic-shift from only the last column by determining how many characters are lexicographically smaller than the current character, and also the rank of the current character at this position. The sum of these two values is the location of the previous cyclic-shift. This function is called the Last-to-first mapping (LF-mapping). Calculating the LF-mapping can be done efficiently by using a rank/select data structure to calculate the rank of all the characters in the BWT, and an array C which contains the number of occurrences of each letter in the BWT.

The LF-mapping will also be useful in the detection algorithm when dynamically updating the suffix array.

Chapter 3

Implementation: LSP server architecture

CCDetect-LSP is integrated into IDEs via the Language server protocol. The tool starts up as an LSP server, which an IDE client can connect to and send/receive messages from. The goal of the LSP server/client interaction is to give users of the tool an overview of all clones as they appear. The LSP server is implemented in Java with the LSP4J library which provides an abstraction layer on top of the protocol which is easier to work with programmatically.

The LSP server shows error-messages (diagnostics) which indicate which section of code a clone covers, they also provide information about the matching clones and code-actions to navigate between them. Whenever source code is edited, the LSP server receives a message containing information about what has changed. The LSP server then generates a new list of clones, and displays those to the client instead.

The following user-stories shows how interaction with the LSP server works.

- A programmer wants to see code clones for a file in their project, the programmer opens the file in their IDE and is displayed diagnostics in the code wherever there are detected clones. The matching code clones are not necessarily in the same file.
- A programmer wants to see all code clones for the current project. The programmer opens the IDEs diagnostic view and will see all code clones detected as diagnostics there. The diagnostic will contain information like where the clone exists, and where the matching clone(s) are.
- A programmer wants to jump to one of the matches of a code clone in their editor. The programmer moves their cursor to the diagnostic and will see a list of the matching code clones. The programmer will select the wanted code clone which will move the cursor to the file and location of the selected code clone. Alternatively, a code-action can be invoked to navigate, if the client does not implement the `DiagnosticRelatedInformation` interface.
- A programmer wants to remove a set of clones by applying the “extract-method” refac-




Figure 3.1: Tool architecture

toring. The programmer performs the necessary refactorings, and will see that all the previously detected clones are gone when the LSP server updates.

Figure 3.1 shows the architecture of the tool. The LSP module communicates with the IDE and delegates the work of handling documents to the document index. Detection of clones is delegated to the detection module. The LSP module receives a list of clones from the detection module which is then converted to LSP diagnostics and code-actions which is finally sent to the client.

3.1 Document index

Upon starting, the LSP server requires indexing of the project for conducting analysis. This involves creating an index and inserting all the relevant documents in the code base. A document contains the content of a file along with extra information such as the file's URI and some information which is useful for the later clone detection. We define the following record for our documents:

```

record Document {
    String uri,
    String content,
    AST ast,

    // Location in fingerprint

```

```

    int start,
    int end

    // Used for incremental updates
    int[] fingerprint,
    boolean open,
    boolean changed,
}

```

Each document in the index primarily consists of the contents, the URI and the AST of the document in its current state. Storing the AST will be useful for the incremental detection algorithm.

There are two things to consider when determining which files should be inserted into the index. First, we are considering only files of a specific file type, since the tool does not allow analysis of multiple programming languages at the same time. Therefore, the index should contain for example only `*.java` files if Java is the language to analyze. Secondly, all files of that file type might not be relevant to consider in the analysis. This could for example be generated code, which generally contains a lot of duplication, but is not practical or necessary to consider as duplicate code, since this is not code which the programmer interact with directly. Therefore, the default behavior is to consider only files of the correct file type, which are checked into Git. The tool supports adding all files in a folder, or all files checked into Git.

When a document is first indexed by the server, the file contents is read from the disk. However, as soon as the programmer opens this file in their IDE, the source of truth for the files content is no longer on the disk, as the programmer is changing the file continuously before writing to the disk. The LSP protocol defines multiple RPC messages which the client sends to the server in order for the server to keep track of which files are opened, and the state of the content of opened files.

Upon opening a file, the client will send a `textDocument/didOpen` message to the server, which contains the URI for the opened file. The index will at this point set the flag `open` for the relevant document and stops reading its contents from disk. Instead, updates to the file are obtained via the `textDocument/didChange` message. This message can provide either the entire content of the file each time the file changes, or it can provide only the changes and the location of the change. Receiving only the changes will be useful for this algorithm when the analysis incrementally reparses the document.

3.2 Displaying and interacting with clones

When detection is finished and the list of clones is ready, the LSP module will convert clones into diagnostics and code-actions which can be interacted with from the client. For diagnostics, each clone object is converted into a diagnostic which has a range and some information about the clone. The diagnostic displays an error in the client which also has information for each matching clone. This is achieved this by attaching `DiagnosticRelatedInformation` to a diagnostic which is defined in the protocol. When all the code clones have been converted to diagnostics, the server sends a `textDocument/publishDiagnostics` which sends all the diagnostics in JSON-RPC format. For code-actions, the process is similar, for each clone pair, we create a code-action which navigates from one to the other, using the `window/showDocument` message. These code-

actions are similarly sent to the client via the `textDocument/codeAction` message, but this message is only invoked when the client specifically requests code actions at a certain location.

Figure 3.2 shows how clones are displayed in the VSCode IDE. Code clones are displayed in-line in the file as an “Information” diagnostic, and when hovered over, the client shows the `DiagnosticRelatedInformation` information of the diagnostic where all the matching clones are displayed and can be clicked to navigate to the corresponding match. For a client which may not implement the `DiagnosticRelatedInformation` message (which was added in LSP version 3.16), invoking the code-action to navigate to the matching clone is another option.

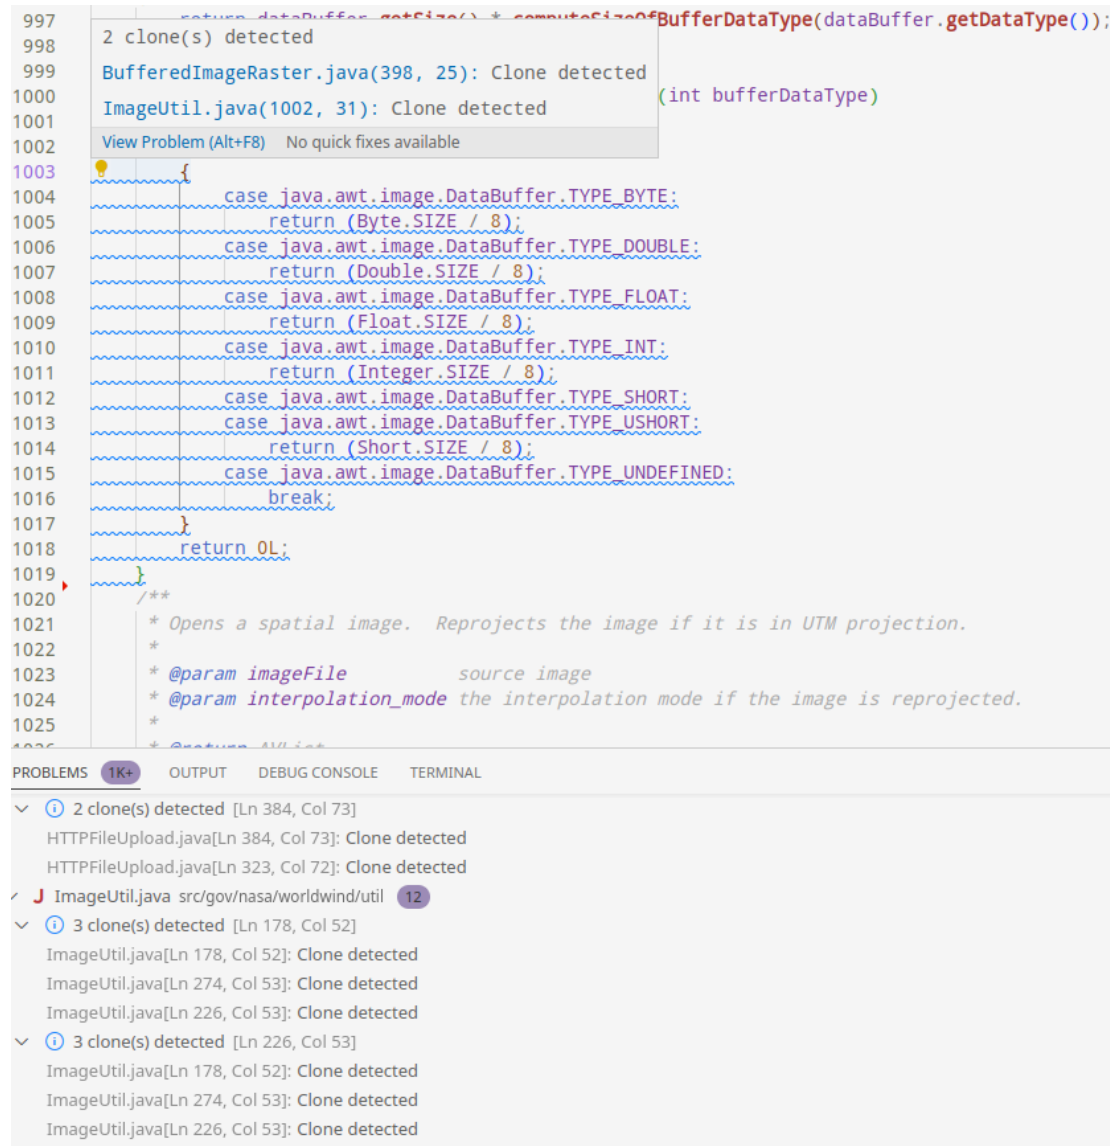


Figure 3.2: Code clones displayed in the VSCode IDE

Chapter 4

Implementation: Initial detection

This chapter discusses the detection module of CCDetect-LSP. It consists of the detection algorithm which takes the document index as input, and outputs a list of code clones. The initial input to the algorithm will be the raw source code of each document in the index, in text format.

The algorithm detects syntactical type-1 code clones, based on a token-threshold. Clones detected are therefore snippets of source-code which has at least N equal tokens, where N is a configurable parameter. There are also two types of clones which are filtered:

- Clones which are completely contained within another larger clone are filtered.
- Clones cannot extend past the end of a fragment.

In broad strokes, the algorithm first selects the relevant parts of source code to detect code clones in (fragment selection), then transforms the selected fragments into a smaller representation (fingerprinting). For the matching, an extended suffix array for the fingerprint is constructed, where the LCP array is used to find long matching instances of source code. Finally, clones are filtered and aggregated into clone classes before they are source-mapped back to the original source-code locations, which the LSP server can display as diagnostics.

4.1 Fragment selection

The first phase of the algorithm involves extracting the relevant fragments of source code which should be considered for detection. A fragment in this case is considered as a section of abstract syntax, meaning that a particular type of node in the AST and the tokens it encompasses should be extracted. Since the algorithm is language agnostic, it is not feasible to have a single algorithm for fragment extraction or to define a separate fragment extraction algorithm for every possible language. Therefore, a parser generator tool is used, which can generate the code for parsing any programming language given a grammar. We use Tree-sitter, a parser generator with incremental parsing and error-recovery capabilities, as well as a query language for finding any type of node in the AST [6]. Tree-sitter queries are flexible queries to extract specific nodes or subtrees. For example, in Java it would be natural to consider only methods. The Tree-sitter query for selecting only the method nodes in a Java programs AST would be:

$$(\text{method_declaration } @\text{method}) \quad (4.1)$$

This Tree-sitter query selects the node with type `method_declaration` and “captures” it with the `@method` name, so that it can be further processed by the program. Readers interested in the details of the query system are referred to the Tree-sitter documentation [6].

Using a tree-sitter query, the algorithm parses the program into an AST and queries the tree for a list of all nodes which matches the query. For each node, all the tokens which it covers (the values in all the leaves in its subtree) are extracted and sent to the next phase.

TODO: Algorithm?

Implementing something similar for another parser/AST could be as simple as traversing the tree until a node of a specific type is found, using the visitor pattern [11, p. 366].

4.2 Fingerprinting

TODO: Høyere abstraksjonsnivå med infinite alphabets

	Token	Fingerprint
	public	2
	int	3
	multiplyByTwo	4
	(5
	param	6
)	7
	{	8
	return	9
	*	10
	2	11
	;	12
	}	13
	addTwo	14
	+	15

```

public class Math() {
    public int multiplyByTwo(int param) {
        return param * 2;
    }

    public int addTwo(int param) {
        return param + 2;
    }
}

```

(a) Source-code

(b) Fingerprint mapping

[2, 3, 4, 5, 3, 6, 7, 8, 9, 6, 10, 11, 1, 2, 3, 14, 5, 3, 6, 7, 8, 9, 6, 15, 11, 1, 0]

(c) Fingerprint, terminals underlined

Figure 4.1: Example fingerprint of Java source-code

The next phase of the algorithm is to transform the extracted tokens into a representation which is less computationally heavy for the matching algorithm. The goal is to reduce the total size of the input which needs to be processed.

Since the algorithm that is used for matching is based on a suffix array, the representation should

be in a format similar to a string, which is the standard input to a suffix array construction algorithm. However, it is not strictly necessary to use strings. The essential property of the input array is that we have a sizeable alphabet where each element is comparable. We will use an array of integers instead of a string as it has a large alphabet in most programming languages.

The algorithm utilizes fingerprinting in order to reduce the size of the representation. Fingerprinting is a technique which involves taking some part of the input and mapping it to a smaller bit string, which uniquely identifies that part. In this case, the algorithm takes each token of the source-code, and maps it to the bit string of an integer. Each unique token value is mapped to unique integer values. Figure 4.1 shows how a sample Java program could theoretically be fingerprinted. The fingerprint mapping starts its “count” at 2 to make space for some terminating characters. Each fragment is terminated by a 1 and the fingerprint is ultimately terminated by a 0. Having these values in the fingerprint will be useful for the matching algorithm where the suffix array is constructed and utilized for detection. The fingerprint of each fragment is stored in the relevant document object.

4.3 Suffix array construction

The next step is to input the fingerprint into a suffix array construction algorithm (SACA), so that the suffix array can be used to find repeating sequences. Recall that the suffix array of a string sorts all the suffixes of the string, as discussed in section 2.7. This makes it simple to find similar regions of text, since similar suffixes will be next to each other in the suffix array. All the fingerprints which are stored in each document object will now be concatenated to be stored in a single integer array (with terminators after each fragment), which the suffix array (SA), inverse suffix array (ISA) and longest-common prefix array (LCP) is computed from.

We have utilized a straight-forward implementation of the “Induced sorting variable-length LMS-substrings” algorithm [23] which computes a suffix array in linear time. The following section will give high-level overview of how the algorithm works.

The algorithm will be assumed to have a string as its input, as this is most common for suffix arrays, and working with strings can be more clear to the reader when considering suffixes. The algorithm is still applicable to our fingerprint, since an array of integers will work similarly to a string when given as input.

The “Induced sorting variable-length LMS-substrings” algorithm (often abbreviated SA-IS) is an algorithm that works by divide and conquering the suffix array and inducing how to sort the suffixes of the string S from a smaller string, S_1 . S_1 consists of the “building blocks” of S and will make it simple to compute the rest of the suffix array of S . First, we will introduce some definitions and theorems used for the algorithm. Input to the algorithm is a string S with length n . Let $\text{suffix}(S, i)$ be the suffix in S starting at position i .

Definition 10 (L-type and S-type suffixes). *A suffix starting at position i in a string S is considered to be L-type if it is lexicographically larger than the next suffix at position $i + 1$. Meaning that $\text{suffix}(S, i) > \text{suffix}(S, i + 1)$. Conversely, a suffix is considered to be S-type if $\text{suffix}(S, i) < \text{suffix}(S, i + 1)$. The sentinel suffix (\$) of S is always S-type.*

Note that two suffixes in the same string cannot be lexicographically equal, therefore all cases are handled by this definition. Determining the type of each suffix can be done in $O(n)$ time

B	A	N	A	N	A	\$
L	<u>S</u>	L	<u>S</u>	L	L	<u>S</u>
0	1	2	3	4	5	6

Table 4.1: Suffix types of $S = \text{BANANA}\$,$ LMS characters underlined

$$\begin{aligned} \text{types}[n] &= S \\ \sum_{i=0}^{n-1} \text{types}[i] &= \begin{cases} S & \text{if } S[i] < S[i+1] \\ L & \text{if } S[i] > S[i+1] \\ \text{types}[i+1] & \text{if } S[i] = S[i+1] \end{cases} \end{aligned}$$

Figure 4.2: Suffix type recurrence

by scanning S from right-to-left and observing the following properties: $\text{suffix}(S, i)$ is L-type if $S[i] > S[i+1]$. Similarly, $\text{suffix}(S, i)$ is S-type if $S[i] < S[i+1]$. If $S[i] = S[i+1]$, then $\text{suffix}(S, i)$ is the same value as $\text{suffix}(S, i+1)$. This is true because if the first character of the current suffix is not equal to first character of the next suffix, we already know the type based on the first character. If the first character is equal, we have effectively transformed the problem to finding the type of the next suffix, since we are now comparing the second character of the current suffix, with the second character of the second suffix. Since we have already computed the type of the next suffix, we can reuse the value. Figure 4.2 shows a recurrence which determines the type of each suffix in a string in $O(n)$ time and 4.1 shows an example.

Definition 11 (LMS character). *An LMS (Left-most S-type) character in a string S is a position i in S such that $S[i]$ is S-type and $S[i-1]$ is L-type. An LMS-suffix is a suffix in S which begins with an LMS character. The final character of S (the sentinel) is always an LMS character and the first character is never an LMS character.*

Definition 12 (LMS-substring). *An LMS-substring in a string S is a substring $S[i..j]$ in S such that $i \neq j$, $S[i]$ and $S[j]$ are LMS characters and there are no other LMS characters between. The sentinel character is also an LMS-substring and is the only LMS-substring of length ≤ 3*

LMS-substrings form "basic-blocks" in the string S . Each LMS-substring is mostly lexicographically decreasing or increasing, which is easier to sort. Table 4.2 shows that in the string BANANA\$ there are 3 S-type suffixes and all of them form LMS-substrings (AN, ANA, \$). Using this notion, we can sort all suffixes recursively using the following theorems:

Theorem 2. *Given sorted LMS-suffixes of S , the rest of the suffix array can be induced in linear time.*

Theorem 3. *There are at most $n/2$ LMS-substrings in a string S of length n .*

We can construct a smaller string S_1 by first sorting the LMS-substrings. Sorting LMS-substrings can be done by first bucket-sorting each LMS-substring by its first character. The buckets are represented by arrays and each LMS-substring is inserted at the end of the correct bucket. Afterwards, L-type suffixes are bucketed by iterating over the buckets, and for each suffix in the bucket, insert the suffix to the left, if it is L-type. Meaning that if we encounter the suffix at

Buckets	\$	A	B	N
Initial	{ -1 }	{ -1, -1, -1 }	{ -1 }	{ -1, -1 }
Slot LMS	{ 6 }	{ -1, 3, 1 }	{ -1 }	{ -1, -1 }
Slot L-type	{ 6 }	{ 5, 3, 1 }	{ 0 }	{ 4, 2 }
Slot S-type	{ 6 }	{ 5, 3, 1 }	{ 0 }	{ 4, 2 }

Mapping	\$	0
	ANA\$	1
	ANA	2

S_1	210
SA of S_1	[2, 1, 0]
Sorted LMS-substrings	[6, 3, 1]

Table 4.2: Building S_1 and sorting LMS-substrings of $S = \text{BANANA}\$$

position 6, the suffix at position 5 is bucketed if it is L-type. Finally, S-type suffixes are bucketed similarly in the same fashion, but the buckets are scanned from right-to-left and suffixes are inserted at the end of the bucket. This final step could possibly change the ordering of the LMS-substrings.

After sorting, each equal LMS-substring is given a unique increasing integer value. Two LMS-substrings are considered equal if they are equal in terms of length, characters and types. S_1 is now built by mapping each LMS-substring to its unique value, and concatenating them in the original ordering. S_1 is now a smaller case which is used in the recursion. The string is at most $n/2$ in size, meaning there will be at most $\log_2(n)$ recursive calls. The recursive call will return the suffix array of the S_1 . Nong et al. proves a theorem which shows that sorting S_1 is equivalent to sorting the LMS suffixes of S . Therefore, the suffix array of S_1 can be mapped to the LMS suffixes of S [23].

The base-case of the recursion is when the suffix array can be computed simply by bucketing each suffix, which happens when every suffix of the string starts with a unique character.

Table 4.2 shows how the LMS-substrings are bucketed and how S_1 is constructed. We see that the final buckets are actually equal to the final SA which we are trying to compute. This is because the LMS-substrings are already sorted in reverse order, which is not true for any arbitrary input. S_1 consists of only unique characters, therefore the suffix array of S_1 is computed by simply bucketing each suffix and returning the array.

When the recursive call returns, the SA of S_1 is used to determine the order which the LMS-substrings should be slotted into the bigger SA. By scanning the smaller SA and mapping those indices back to the indices of the original LMS-substrings, we get a sorted ordering of the LMS-substrings. We then scan the sorted LMS-substrings from right-to-left and slot each LMS-substring at the end of its bucket, the rest of the L-type and S-type suffixes will be slotted correctly afterwards. For $\text{BANANA}\$$ this will be the exact same process as in table 4.2, since the LMS-substrings were already inserted in reverse ordering.

There will be $O(\log_2(n))$ recursive call, where each recursive call takes $O(n)$ time, with n halving in each call. Therefore, the recurrence will have the complexity of:

$$T(n) = \begin{cases} O(n) & \text{if base-case.} \\ T(n/2) + O(n) & \end{cases} = O(n)$$

Building ISA and LCP arrays

Computing the ISA after constructing the SA is simple. Since the ISA is simply the inverse of SA, it can be constructed in linear time with a single loop as seen in Algorithm 3.

Algorithm ComputeISA(SA)

```

n ← SA.len
ISA ← array of size n
for i from 0 to n do
  | ISA[SA[i]] ← i
end
return ISA

```

Algorithm 3: Compute ISA from SA

Computing the LCP in linear time is more complicated and requires some insight about which order to insert LCP values. We will also add one extra restriction to the LCP values, being that the LCP values cannot match past a 1, which were the terminal value between fragments. This restriction will be useful when we want to extract clones using the LCP array. The algorithm to compute the LCP is shown in Algorithm 4. The intuition for this algorithm is that if a suffix at position i has an LCP value l describing the common-prefix between it and some other suffix at position j , then the LCP value of the suffix at position $i + 1$ is at least $l - 1$, since the suffix at $i + 1$ and $j + 1$ is the same suffix as the suffixes at position i and j , with the first character cut off. Therefore, they share at least $l - 1$ characters, and the algorithm can start comparing the characters at that offset.

Algorithm ComputeLCP(S, SA, ISA)

```

n ← SA.len
LCP ← array of size n
lcpLen ← 0
for i from 0 to n - 1 do
  | r ← ISA[i]
  | prevSuffix ← SA[r - 1]
  | while S[i + lcpLen] = S[prevSuffix + lcpLen] and S[i + lcpLen] ≠ 1 do
  | | lcpLen ← lcpLen + 1
  | end
  | LCP[r] ← lcpLen
  | lcpLen ← Max(0, lcpLen - 1)
end
return LCP

```

Algorithm 4: Compute LCP from input string S, SA, and ISA

Index	Suffix						Minimum LCP
1	A	N	A	N	A	\$	0
3	A	N	A	\$			0
2	N	A	N	A	\$		2
4	N	A	\$				2

Table 4.3: Minimum common LCP values between suffixes for $S = \text{BANANA}\$$

Algorithm SimpleCloneExtraction(S, ISA, LCP)

```

 $n \leftarrow S.len$ 
clones  $\leftarrow$  list
for  $i$  from 0 to  $n - 1$  do
    if  $ISA[i] = 0$  then
        | continue
    end

    if  $LCP[ISA[i]] \geq THRESHOLD$  then
        | Insert(clones,  $i$ ) // Adds  $i$  to the clone-list
    end
end
return clones

```

Algorithm 5: Extract clones indices in a string S

4.4 Clone extraction

With the extended suffix array computed, we can now consider which substrings (prefixes of suffixes) we want to extract as potential code clones. In this phase the indices of the fingerprint which we consider to be code clones are extracted, which will be mapped back to the original source code in the next phase.

A straightforward solution is to extract every suffix which has an LCP value which is greater than the token threshold. The algorithm is a loop over S , using ISA to find the corresponding LCP value. This finds the clone indices, as shown in Algorithm 5.

However, this algorithm will return a lot of contained clones. A contained clone is a clone where all the tokens of the clone is a part of another, larger clone. The algorithm will give a lot of contained clones, for example in the case where there is a suffix with an LCP value of 100, the next suffix will have the LCP value of at least 99 and likely matches with the same code clone as the previous suffix, but with an offset of 1 token. For any large code clone, there will therefore be many smaller clones which are completely contained within it, but these clones are also likely to match with another clone which is also contained within the larger clones match. Since the code clones point at mostly the same area, the contained clones are not very useful, and should not be considered. We extend our clone extraction algorithm to account for this, by using the following theorem:

Theorem 4. *The LCP of a suffix at position i is completely contained in the LCP of the previous suffix at position $i - 1$ if the LCP value of the suffix at position $i - 1$ is greater than the LCP value of the suffix at position i . Meaning $LCP[ISA[i]] < LCP[ISA[i - 1]]$.*

Algorithm CloneExtraction(S, ISA, LCP)

```

n ← S.len
clones ← List
for i from 0 to n - 1 do
  if ISA[i] = 0 then
    | continue
  end

  if LCP[ISA[i]] ≥ THRESHOLD then
    Insert(clones, i) // Adds i to the clone-list

    while i + 1 < n and LCP[ISA[i + 1]] < LCP[ISA[i]] do
      | i ← i + 1
    end
  end
end
return clones

```

Algorithm 6: Extract clones indices in a string S , ignoring contained clones

Algorithm 6 adds a while-loop to the clone extraction algorithm which skips over suffixes which are contained according to the theorem. Note that this algorithm doesn't disallow contained clones entirely, but any clone which is a shorter version of another clone pointing to the same match, will be skipped. Also note that overlapping clones, meaning two clones which share tokens, but where neither contains the other, will not be skipped.

Finally, we have a list of indices in the fingerprint which is considered to be code clones.

4.5 Source-mapping

With the clone indices in the fingerprint computed, we are almost finished. The final step is to map the clone indices back to the original source code. In order to correctly identify where a clone is located, we need to know which file the clone is located in, the range of the source code in that file the code clone covers, and the other matching code clones.

To accomplish this, each document in the index needs to store the range of each of its tokens and keep track of which portion of the fingerprint consists of the documents tokens. This is done by storing two integer variables, each storing the start position and end position that the document has in the fingerprint.

To determine which document a fingerprint position corresponds to, we can perform a binary search on the list of the documents, which is sorted based on the start position of the documents fingerprint. The goal is to find the document D where the fingerprint position i is

$$D.start \leq i \leq D.end$$

Once the correct document has been found, we simply have to look up the correct range which

Algorithm SourceMap(*documents*, *i*)

```

left ← 0
right ← documents.len - 1

while left ≤ right do
  mid ← (left + right)/2
  if documents[mid].end < i then
    | left = mid + 1
  else if documents[mid].start > i then
    | right = mid - 1
  else
    | D ← documents[mid]
    | range ← D.ranges[i - D.end]
    | return (D.uri, range)
  end
end
end

```

Algorithm 7: Get source-map for a position *i* in the fingerprint

the document stores. The index of this range is $i - D.start$.

Algorithm 7 shows this algorithm which outputs the URI for the document and the source code range of the token at position *i*.

This algorithm only shows how to look up the position of a single token. Since a code clone is a range between two tokens, we have to look up the position at index *i* extracted in the previous phase, and the position where the clone ends, which is index $i + \text{LCP}[\text{ISA}[i]]$. The range of the code clone is therefore the combination of the starting range of the first token (at position *i*) and the ending range of the second token (at position $i + \text{LCP}[\text{ISA}[i]]$):

Aggregating clones

With this algorithm to get the clone locations, the next step is to make sure that matching code clones are collected into buckets of clone classes. Remember that the LCP array only gives us the longest match between two suffixes, but it is naturally possible to have more than two clones of the same code snippet. This case happens when multiple consecutive indices in the SA are considered to be clones. Since we only look for type-1 clones, the transitivity property holds, meaning that if

$$SA[i] \xrightarrow{\text{clone}} SA[i + 1] \xrightarrow{\text{clone}} SA[i + 2]$$

then clones at position $SA[i]$, $SA[i + 1]$ and $SA[i + 2]$ are all clones of each other. This should be achieved by making sure that every new clone-pair discovered adds previously detected clones to their clone sets, and previously discovered clones also add new clones to theirs. Figure 4.3 shows how a new match is found in two previously disjoint clone classes, and the resulting aggregated clone class.

This is achieved in algorithm 8 where we build a clone-map, where the key is the index that a




Figure 4.3: Clone class aggregation when a new match is found

clone starts at in the fingerprint. For each clone index i which was extracted in the last phase, we get the corresponding match index j ($SA[ISA[i] - 1]$), and for both i and j we look in the clone-map if there already is a clone at that position, or a new clone object is created and put in the map. Then, in order to aggregate the previously discovered clones and the new clone together, the set of matching clones is unioned between the two clones. In this way, all the previous existing clones of i is added as a clone in j and vice versa. Every clone in the two clone classes will then receive the same set of code clones. The **UnionCloneClass** function unions all the clone sets in both clone classes and then adds a match from all clones in one clone class to all clones in the other.

Finally, we have a list of every code clone in the code base, which is then sent to the LSP module of the tool, which handles the displaying of code clones to the client as shown in chapter 3.

Algorithm `GetCloneMap(index, cloneIndices)`

```
n ← cloneIndices.len
cloneMap ← Empty map with type: int → CodeClone

for i from 0 to n − 1 do
  firstIndex ← cloneIndices[i]
  secondIndex ← cloneIndices[i]
  size ← LCP[ISA[firstIndex]] − 1

  // Use SourceMap function to get ranges and documents of clones
  firstRange ← Get source between firstIndex and (firstIndex + size)
  secondRange ← Get source between secondIndex and (secondIndex + size)

  firstDocument ← firstRange.document
  secondDocument ← secondRange.document

  // Build clone objects if they don't exist
  if firstIndex not in cloneMap then
    firstClone ← new CodeClone(firstDocument.uri, firstRange)
    Put(cloneMap, firstIndex, firstClone) // Put clone with firstIndex as key
  end
  if secondIndex not in cloneMap then
    secondClone ← new CodeClone(secondDocument.uri, secondRange)
    Put(cloneMap, i, secondClone) // Put clone with secondIndex as key
  end

  // Union clone classes
  UnionCloneClass(Get(cloneMap, firstIndex), Get(cloneMap, secondIndex))
end
return cloneMap
```

Algorithm 8: Build clone-map given the clone indices in the fingerprint

Chapter 5

Implementation: Incremental detection

The following chapter will present the algorithm which efficiently updates the list of clones, without having to rebuild the different structures from scratch. Given an edit to a file in the project, we will be able to update the document index, fingerprints, suffix array and list of clones faster than the initial detection.

An incremental update is run whenever a document is changed. The document index is signaled of a change either when a file is saved, or on any keystroke. This is configurable by the client. When the document index is changed, an incremental update of the clones is run, and the clone-list is updated.

5.1 Affordable operations

Before looking at the approach, it is useful to determine the time cost associated with different operations. If we can for example afford to iterate over the contents of a single file, that will be useful for our algorithm. Table 5.1 shows which operations are affordable or not for an incremental update.

5.2 Updating the document index

The first step of an incremental update is to update the document index. We will also look at how we can reduce memory usage of the index without a loss in terms of the time complexity of the updates.

As shown in the document interface, each document stores its own content, AST and fingerprint. It is not strictly necessary to store either the content or the AST in memory all the time, as it is likely that only a handful of files are open in the IDE at once. Therefore, in the initial detection, we can free the memory of the file content and AST for each document after the fingerprint has been computed. However, if a file is opened in the IDE, the file can now be changed, so we should facilitate efficient updates for these files only. When a file is opened, the file content

Description	Complexity	Affordable	Explanation
Content of all files	$O(\text{code base})$	No	Iterating over the entire code base will be the same complexity as the initial detection, therefore this operation is too slow.
Content of a file	$O(\text{file})$	Yes	Iterating over the content of a single file is likely a very small percentage of the entire code base.
Parsing a file	$O(\text{file})$	No	While the complexity of parsing a single file is still linear in the size of the file, parsing a large file from scratch can take a significant amount of time in practice.
Incrementally reparsing a file	$O(\text{edit})$	Yes	Re-using the AST of a file in order to speed up the parsing of the same file with after an edit is significantly faster than parsing the entire file.
Document index	$O(\text{documents})$	Yes	The number of files in a code base is likely many orders of magnitude smaller than the size of the code base itself.
Clones	$O(\text{clones})$	Yes	The number of clones in the code base and the area they cover is likely a very small portion of the code base itself. itself.

Table 5.1: Affordable operations for incremental updates

should be read from the disk and updated via the `textDocument/didChange` messages sent from the client. It is also important to keep the AST of the opened file in memory in order to facilitate incremental reparsing of the opened files.

When a file is opened, the LSP client sends a `textDocument/didOpen` message to the server, which finds the relevant document D in the index, and sets the following fields:

$$\begin{aligned}
D.open &= \mathbf{True} \\
D.AST &= \text{Parse}(D) \\
D.content &= \text{Read}(D.uri)
\end{aligned}$$

After the document fields have been set, the document is ready to receive updates. When the LSP client sends a `textDocument/didChange` message, the message consists of the URI of the edited file, the range of the content which has changed, and the content which has potentially been inserted. This range is then used in a tree-sitter incremental reparse of the file content. After this reparse, we have efficiently updated a documents content and AST. After this update, we also set $D.changed = \mathbf{True}$

5.3 Updating fingerprints

With the updated AST for all documents, we can update the fingerprint of all documents which have been changed. For each document D where $D.changed = \mathbf{True}$, the fingerprint for D may

have changed. Calculating the fingerprint is the same process as in the initial detection, where we first query the AST for all nodes of a certain type, then for each matched node N , we extract and fingerprint all the tokens which N covers, using the same fingerprint mapping as was used for the initial detection.

An additional change we have to consider when incrementally updating fingerprints is that for a document D , $D.start$ and $D.end$ which corresponds to the range which D covers in the fingerprint, may have changed. Also, any document D_1 where $D_1.start > D.start$ could also have its range changed. This is solved while updating each documents fingerprint by counting the number of tokens in each document after updating, and setting the appropriate **start** and **end** fields.

TODO: Algorithm for updating document index here

TODO: Figure which displays three documents, with old and new fingerprint

5.4 Computing edit operations

Now that the fingerprint has been updated, we could build the suffix array from scratch and already see a substantial improvement in performance. The major bottleneck of the initial detection is to parse and fingerprint the entire code base. However, the updating of the suffix array can and should also be updated incrementally to further improve efficiency.

The input to the dynamic suffix array algorithm is a set of edit operations. Therefore, we need to know what exactly has changed in the fingerprint before we can update the suffix array. We need to determine what edit operations have happened. Edit operations are either deleting, inserting or substituting a section of the code.

There are two approaches we could take to determine the edit operations. The first is to look at the ranges that the LSP client sends with each `textDocument/didChange` message and determine which tokens in the fingerprint have been affected by this message. However, this approach tightly couples the algorithm to LSP and the scenario where we know the exact ranges of each change. Also, we might do unnecessary amounts of operations if we do multiple edits, since some operations could cancel each other out, for example by inserting and then deleting some text.

The other approach is to determine the changes of the fingerprint via an edit distance algorithm. An edit distance algorithm is an algorithm which calculates the distance between two strings S_1 and S_2 . Distance between two strings is the minimum number of edit operations (insert, delete, substitute) which is required to transform S_1 into S_2 . Many of the algorithms which calculates the edit distance also allows computing what the operations are.

The classic algorithm for calculating edit distance operations is attributed to Wagner and Fischer [32]. The input to the algorithm is two strings S_1 and S_2 of length n and m . The output will be the set of operations needed to turn S_1 into S_2 . This algorithm is based on dynamic programming where a matrix M is filled from top to bottom and then the operations are inferred from M . Algorithm 9 shows how the edit distance matrix is filled and 5.2 shows an example matrix.

Each index i, j in M contains the edit distance value between the substrings $S_1[0..i]$ and $S_2[0..j]$. The values in M is calculated by determining what is the cheapest operation to do at a certain

$$\begin{aligned}
\sum_{i=0}^n M[i][0] &= i \\
\sum_{j=0}^m M[0][j] &= j \\
M[i][j] &= \begin{cases} M[i-1][j-1] & \text{if } S_1[i] = S_2[j] \\ 1 + \text{Min}(M[i-1][j], M[i][j-1], M[i-1][j-1]) & \text{otherwise} \end{cases}
\end{aligned}$$

Figure 5.1: Edit distance recurrence

location to make the substrings equal. This can be determined by looking at the three surrounding indices in M : $M[i-1][j-1]$, $M[i-1][j]$ and $M[i][j-1]$. Each of these indices equate to deleting, inserting or substituting a character in S_1 . The recurrence in figure 5.1 describes the algorithm.

The edit operations can then be inferred from M by backtracking from the bottom-right index, to the top-left, giving us the edit operations in reverse. At each position i, j we choose either of the 3 surrounding indices, the same indices which were used to determine the value originally. Choosing the left index $(i, j-1)$ equates to inserting the character $S_2[j-1]$ at position $i-1$. Choosing the top index $(i-1, j)$ equates to deleting the character $S_1[i-1]$. Choosing the top-left index $(i-1, j-1)$ equates to substituting $S_1[i-1]$ with $S_2[j-1]$. If these characters are already equal, the operation can be ignored. For example in table 5.2, the first operation is to substitute R with D at position 0. Afterwards P and U is deleted at position 2. Then the B which is now at position 2 is substituted by M. This continues with more substitutions until we finally have DEMOCRAT.

TODO: Now collect operations together and explain memory issue + Hirschbergs

Aggregating edit operations

In the next phase we will feed the edit operations into an algorithm which dynamically updates our suffix array based on those operations. However, this algorithm will be more efficient if the operations are combined to singular inserts, deletes or substitutes of strings more than one character. The edit distance algorithm outputs only single character operations, meaning we insert, delete or substitute a single character at a time. We therefore want a way to combine these operations to “larger” operations.

We define an EditOperation with the following record:

```
record EditOperation {
    OperationType type,
    char[] chars,
    int position
}
```

where type is either an insert, delete or substitute.

Algorithm WagnerFischerEditDistance(S_1, S_2)

```

n ← Len( $S_1$ )
m ← Len( $S_2$ )
matrix ← new array[n + 1][m + 1]

for  $i$  from 0 to  $n$  do
  | matrix[i][0] =  $i$ 
end

for  $i$  from 0 to  $m$  do
  | matrix[0][i] =  $i$ 
end

for  $i$  from 1 to  $n$  do
  for  $j$  from 1 to  $m$  do
    if  $S_1[i - 1] = S_2[j - 1]$  then
      | matrix[i][j] = matrix[i - 1][j - 1]
    end
    else
      delete ← matrix[i - 1][j]
      insert ← matrix[i][j - 1]
      substitute ← matrix[i - 1][j - 1]

      matrix[i][j] = Min(insert, delete, substitute) + 1
    end
  end
end
return matrix

```

Algorithm 9: Fill edit distance matrix using Wagner-Fischer algorithm

		D	E	M	O	C	R	A	T
	0	1	2	3	4	5	6	7	8
R	1	1	2	3	4	5	5	6	7
E	2	2	1	2	3	4	5	6	7
P	3	3	2	2	3	4	5	6	7
U	4	4	3	3	3	4	5	6	7
B	5	5	4	4	4	4	5	6	7
L	6	6	5	5	5	5	5	6	7
I	7	7	6	6	6	6	6	6	7
C	8	8	7	7	7	6	7	7	7
A	9	9	8	8	8	7	7	7	8
N	10	10	9	9	9	8	8	8	8

Table 5.2: Edit distance matrix for REPUBLICAN → DEMOCRAT

One way to do this is to combine operations of the same type which are sequenced in the matrix. For example in table 5.2, we have two consecutive delete operations, where P and U is deleted at position 2. These two operations could be combined to a single delete operation of two characters.

The idea for the algorithm which computes the edit operations is to traverse the optimal matrix path backwards and for each operation we either append to the current operation if possible, or start a new operation. We can continue the current operation if the next operation is of the same type, and the next operation has the same position as the current operation. For example in table 5.2, we have two delete operations in a sequence at position 1 where P and U is deleted. The first operation we encounter is the deletion of U. At this point we create a new **EditOperation** with position 1 and U in the **chars** array. The next operation is also a delete operation at position 1, so we add the P to the list of characters for the operation to delete. The next operation is a substitute at position 0, so we cannot continue the delete operation, and a new substitute operation is created instead. Similarly, at position 2 to 5, we substitute **BLIC** with **MOCR**. This operation is computed in a backwards fashion similarly to how we did the deletion, except that the position of the operation is decremented for each character we add to it.

Optimizing memory usage and operations

A problem with the above solution is the memory usage of the matrix. It is not feasible to input the entire old and new fingerprint into an edit distance algorithm, as the full fingerprint can have many million symbols, and the old and new fingerprint is likely approximately the same size. This would require a matrix which is too big in memory. We will use a few techniques to reduce the memory usage of this algorithm without compromising on the time complexity.

The first technique is to not input the whole fingerprint. We can drastically reduce the size of the input by only comparing the old and new fingerprint of the document *D* which has been edited. *D* stores its previous and current fingerprint, and whenever it is edited, we can compute the edit operations of *D*, with its previous and current fingerprint as input. However, the position of each edit operation of *D* will not have the correct position relative to the entire fingerprint, so *D.start* is added to the position of each edit operation to correct this.

Another optimization we can do to reduce the size of the matrix is to remove the “trivial” part at each end of our matrix. If we compare two strings which have many similar characters at the beginning or end of our string, we know that these will not equate to any edit operations in the matrix, as they will simply be diagonal moves (substitutes) where the characters are already equal. In table 5.3 we see the edit matrix for the input **FASCINATING** and **FINISHING**. The two words share the common prefix **F** and the common suffix **ING**. If we examine the edit distance values, we see that the highlighted green matrix starts at 0 in the top left, and ends with 6 in the top right, which is exactly the same values as in the full matrix. Also, we see that there are no edit operations being added outside the green matrix. Using this knowledge, we can see that we only need to compare the strings **ASCINAT** and **INISH**, which would give us the exact same edit distance. We can also get the exact same edit operations, as long as we account for the starting offset of the original string, so that the operations have the correct positions. Algorithm 10 shows how two input strings can be minimized for usage in the edit distance algorithm. After getting the edit operations from the algorithm, the **startOffset** is added to the position of each operation to account for the offset of the minimized matrix.

The two optimizations drastically reduce the memory usage of the edit distance algorithm, but in cases where a very large file is edited, and the beginning and end of the old and new fingerprint

		F	I	N	I	S	H	I	N	G
	0	1	2	3	4	5	6	7	8	9
F	1	0	1	2	3	4	5	6	7	8
A	2	1	1	2	3	4	5	6	7	8
S	3	2	2	2	3	3	4	5	6	7
C	4	3	3	3	3	4	4	5	6	7
I	5	4	3	4	3	4	5	4	5	6
N	6	5	4	3	4	4	5	5	4	5
A	7	6	5	4	4	5	5	6	5	5
T	8	7	6	5	5	5	6	6	6	6
I	9	8	7	6	5	6	6	6	7	7
N	10	9	8	7	6	6	7	7	6	7
G	11	10	9	8	7	7	7	8	7	6

Table 5.3: Edit distance matrix for FASCINATING \rightarrow FINISHING. Blue is optimal path, green is minimized matrix

Algorithm EditDistanceMinimizeStrings(S_1, S_2)

```

for  $i$  from 0 to  $\text{Min}(\text{Len}(S_1), \text{Len}(S_2))$  do
  if  $S_1[i] \neq S_2[i]$  then
    break
  end
end
startOffset  $\leftarrow i$ 

s1End  $\leftarrow \text{Len}(S_1) - 1$ 
s2End  $\leftarrow \text{Len}(S_2) - 1$ 
while s1End  $\geq$  startOffset and s2End  $\geq$  startOffset and  $S_1[s1End] = S_2[s2End]$  do
  s1End  $\leftarrow$  s1End - 1
  s2End  $\leftarrow$  s2End - 1
end

miniS1  $\leftarrow S_1[\text{startOffset} \dots s1End]$ 
miniS2  $\leftarrow S_2[\text{startOffset} \dots s2End]$ 

return (miniS1, miniS2, startOffset)

```

Algorithm 10: Minimize strings for edit distance algorithm

		D	E	M	O	C	R	A	T
R									
E									
P	3	3	2	2	3	4	5	6	7
U	4	4	3	3	3	4	5	6	7
B									
L									
I									
C									
A									
N									

Table 5.4: Edit distance matrix with minimal memory usage

do not match, we can still encounter instances of the matrix being too large to fit in memory. The problem is that the matrix size has a polynomial growth in terms of the fingerprint size. This is because for the old fingerprint of size n and the new fingerprint of size m , we have that $n \approx m$, which requires an $n * m$ size matrix to calculate the edit operations. For example if a Java file contains 3000 lines, the number of tokens can exceed 10000, which would require approximately a $10000 * 10000$ size matrix, which is approaching a memory usage which too large.

A solution to this problem is to reduce the required memory from the polinomial $O(n * m)$ memory usage, to a linear $O(n)$ memory usage. A stepping stone towards such a solution is an observation attributed to Ukkonen, which reduces the required memory to linear growth in the size of either of the strings[30]. The observation shows that in order to compute the next row/column of the edit distance matrix, we only need the previous row/column. For example in table 5.4, the fifth row has been computed using only the fourth row. This is done by first computing the left-most index of the fifth row, which is always one more than the previous row. Afterwards, we can compute the other elements of the row from left to right, with the same recurrence, since we always know the above, left, and top-left elements of the current index. This can be implemented as two arrays, for the previous and current row. The optimization can even be implemented using only a single array, but this does not change the complexity of the memory usage.

This change allows us to compute the edit distance in linear space, but now the problem is how to find the actual edit operation. This is not possible, because we don't have the entire matrix available to traverse anymore. However, this can be solved using Hirschberg's algorithm[13]. Hirschberg's algorithm is an algorithm which can compute the edit operations of two strings in the same time complexity as the Wagner-Fischer algorithm, but uses only linear space.

The first observation we need to make for this algorithm is that there is at minimum one edit operation on each row/column of the edit distance matrix. This is intuitive, because in order to "travel" from the top-left to the bottom-right of the matrix, we need to travel from the top to the bottom, visiting all the rows, and from the left to the right, visiting all the columns. Hirschberg's algorithm uses this to compute one position of the optimal path at a time, which it finds in the middle row between two already known positions of the path.

Initially, we know that the top-left index, and the bottom-right index of the matrix are guaranteed

to be part of the optimal path, since those positions are the starting and ending point of the path. The next position to find is in the middle row of the matrix. Determining which position in the middle row should be selected is done by performing the edit distance recurrence twice, once from the beginning to the middle row in the same fashion as the original algorithm, and once in reverse, from the bottom-right to the middle row. Both of the recurrences can run in linear space, because we only need to store two rows at a time. Summing the two results for the middle row gives us an array which can intuitively be understood as how much the minimal path which goes through the corresponding position in the row costs. Since we know that there is at least one of the positions in this row which has to be part of the optimal path, the minimum value in this array corresponds to one position which has to be part of the optimal path. Once we know the middle position which is part of the optimal path, we can recursively call the function twice, once with the top-left and middle as input, and once with the middle and bottom-right as input. The base-case of the recursion is when the size of the matrix is linear in the size of either of the strings, in which case we call the Wagner-Fischer edit distance recurrence with a linear space complexity.

Each position which is a part of the optimal path is stored as they are found, and can then be traversed backwards to determine the edit operations similarly to traversing the matrix. For example if we iterate over the list of positions backwards, we will know that if the first position is at x, y and the second position is at $x - 1, y$, that corresponds to a delete operation, just as in the matrix-based algorithm.

Using Hirschberg's algorithm, we are now able to compute the edit operations for very large files without memory usage issues.

5.5 Dynamic suffix arrays

5.6 Storing old clones

		F	I	N	I	S	H	I	N	G
F										
A										
S										
C										
I	10	8	6	7	6	7	8	8	10	12
N										
A										
T										
I										
N										
G										

		F	I	N	I	S	H	I	N	G
F										
A	5	3	3	4	6					
S										
C										
I										
N										
A										
T				5	4	3	4	6	8	
I										
N										
G										

		F	I	N	I	S	H	I	N	G
F	0	1	2							
A	1	0	1							
S	2	1	1							
C			2	2	4					
I										
N					3	3	4			
A										
T										
I						2	0	2	4	
N										
G										

		F	I	N	I	S	H	I	N	G
F										
A										
S				0	1					
C				1	1					
I				2	1					
N				0	1	2				
A				1	1	2				
T				2	2	2				
I							0	1	2	
N							1	0	1	
G							2	1	0	

Table 5.5: Hirschberg's algorithm

Chapter 6

Evaluation

We will evaluate this tool based on different criteria, which combined will provide a basis for evaluating the tool as a whole.

Since the tool is focused on efficient detection and management of code clones, real-time performance of the tool will be a high priority in its evaluation. The tool will implement different techniques of detecting and merging clones. These will be empirically compared against each other. The tool will also be evaluated against existing tools empirically. We will utilize BigCloneBench [28] to evaluate detection techniques, by running our detection techniques in a standalone mode. We will distinguish between initial detection and incremental detection when evaluating.

The tool will also be evaluated in how well it fits into the software development cycle. Can we determine if this tool is an effective way to detect a clone early in its lifecycle so that they can be removed before it manifests in the source code? In relation to this, we will evaluate if LSP is a suitable tool for use in clone management and refactoring in general. Can LSP provide all the features one would want in a modern analysis tool? What is missing, and how could the LSP protocol be extended in order to facilitate this? We believe that if LSP is an appropriate tool to use for clone management, LSP will also be an appropriate tool for static analysis tools in general.

6.1 Big-O analysis of incremental detection

6.2 Performance comparison

TODO: Flame graphs?

6.3 Comparison with iClones

6.4 Verifying clones with BigCloneBench

Chapter 7

Discussion

- 7.1 Speed of incremental updates vs linear-time SACA
- 7.2 Comparison with iClones
- 7.3 Chosen clones
- 7.4 Usability while programming

Chapter 8

Conclusion and future work

8.1 Future work

Compressing data structures

Optimal edit operations

Delaying LCP array updates

Refactoring of clones

8.2 Conclusion

Chapter 9

Appendix

Algorithm WagnerFischerOperations(S_1, S_2, matrix)

```

operations  $\leftarrow$  Empty list
currentOperation  $\leftarrow$  None operation
lastOperationIndex  $\leftarrow -1$ 
 $x \leftarrow \text{Len}(\text{matrix}) - 1$ 
 $y \leftarrow \text{Len}(\text{matrix}[x]) - 1$ 

while  $x > 0$  and  $y > 0$  do
    subValue  $\leftarrow \text{matrix}[x - 1][y - 1]$ 
    delValue  $\leftarrow \text{matrix}[x - 1][y]$ 
    insValue  $\leftarrow \text{matrix}[x][y - 1]$ 
    if  $\text{subValue} \leq \text{delValue}$  and  $\text{subValue} \leq \text{insValue}$  and then
        if  $\text{subValue} \neq \text{matrix}[x][y]$  then
            if currentOperation is not a substitute or lastOperationIndex  $\neq x$  then
                currentOperation  $\leftarrow$  New substitute operation with position  $x - 1$ 
                Insert(operations, 0, currentOperation) // Add operation at index 0
            end
            else
                currentOperation.position  $\leftarrow$  currentOperation.position  $- 1$ 
            end
        end
        Add  $S_2[y - 1]$  to beginning of currentOperation.chars
        lastOperationIndex  $\leftarrow x - 1$ 
         $x \leftarrow x - 1$ 
         $y \leftarrow y - 1$ 
    end
    else if  $\text{insValue} \leq \text{delValue}$  then
        if currentOperation is not an insert or lastOperationIndex  $\neq x$  then
            currentOperation  $\leftarrow$  New delete operation with position  $y - 1$ 
            Insert(operations, 0, currentOperation) // Add operation at index 0
        end
        Add  $S_1[x - 1]$  to beginning of currentOperation.chars
        lastOperationIndex  $\leftarrow x - 1$ 
         $y \leftarrow y - 1$ 
    end
    else
        if currentOperation is not a delete or lastOperationIndex  $\neq x$  then
            currentOperation  $\leftarrow$  New insert operation with position  $x - 1$ 
            Insert(operations, 0, currentOperation) // Add operation at index 0
        end
        Add  $S_1[x - 1]$  to beginning of currentOperation.chars
        lastOperationIndex  $\leftarrow x - 1$ 
         $x \leftarrow x - 1$ 
    end
end
return operations

```

Algorithm 11: Compute edit operations from Wagner-Fischer matrix

Bibliography

- [1] Mohamed Ibrahim Abouelhoda, Stefan Kurtz, and Enno Ohlebusch. “Replacing suffix trees with enhanced suffix arrays”. In: *Journal of Discrete Algorithms* 2.1 (2004). The 9th International Symposium on String Processing and Information Retrieval, pp. 53–86. ISSN: 1570-8667. DOI: [https://doi.org/10.1016/S1570-8667\(03\)00065-0](https://doi.org/10.1016/S1570-8667(03)00065-0). URL: <https://www.sciencedirect.com/science/article/pii/S1570866703000650>.
- [2] Raihan Al-Ekram et al. “Cloning by accident: an empirical study of source code cloning across software systems”. In: *2005 International Symposium on Empirical Software Engineering (ISESE 2005), 17-18 November 2005, Noosa Heads, Australia*. IEEE Computer Society, 2005, pp. 376–385. DOI: 10.1109/ISESE.2005.1541846. URL: <https://doi.org/10.1109/ISESE.2005.1541846>.
- [3] Paris Avgeriou et al. “Managing Technical Debt in Software Engineering (Dagstuhl Seminar 16162)”. In: *Dagstuhl Reports* 6.4 (2016), pp. 110–138. DOI: 10.4230/DagRep.6.4.110. URL: <https://doi.org/10.4230/DagRep.6.4.110>.
- [4] Brenda S. Baker and Raffaele Giancarlo. “Sparse Dynamic Programming for Longest Common Subsequence from Fragments”. In: *Journal of Algorithms* 42.2 (2002), pp. 231–254. ISSN: 0196-6774. DOI: <https://doi.org/10.1006/jagm.2002.1214>. URL: <https://www.sciencedirect.com/science/article/pii/S0196677402912149>.
- [5] Ira Baxter et al. “Clone Detection Using Abstract Syntax Trees.” In: vol. 368-377. Jan. 1998, pp. 368–377. DOI: 10.1109/ICSM.1998.738528.
- [6] Max Brunsfeld. *Tree-sitter*. <https://tree-sitter.github.io/tree-sitter/>. Accessed: 2022-04-16.
- [7] Michael Burrows and David J. Wheeler. “A Block-sorting Lossless Data Compression Algorithm”. In: 1994.
- [8] Diego Cedrim et al. “Understanding the impact of refactoring on smells: a longitudinal study of 23 software projects”. In: *Proceedings of the 2017 11th Joint Meeting on Foundations of Software Engineering, ESEC/FSE 2017*. Ed. by Eric Bodden et al. ACM, 2017, pp. 465–475. DOI: 10.1145/3106237.3106259. URL: <https://doi.org/10.1145/3106237.3106259>.
- [9] P.B. Crosby. *Quality is Free: The Art of Making Quality Certain*. New American Library, 1980. ISBN: 9780451624680. URL: <https://books.google.no/books?id=3TMQt73LDooC>.
- [10] Martin Fowler. *Refactoring - Improving the Design of Existing Code*. Addison Wesley object technology series. Addison-Wesley, 1999. ISBN: 978-0-201-48567-7. URL: <http://martinfowler.com/books/refactoring.html>.

- [11] Erich Gamma et al. *Design Patterns: Elements of Reusable Object-Oriented Software*. 1st ed. Addison-Wesley Professional, 1994. ISBN: 0201633612. URL: http://www.amazon.com/Design-Patterns-Elements-Reusable-Object-Oriented/dp/0201633612/ref=ntt_at_ep_dpi_1.
- [12] Nils Göde and Rainer Koschke. “Incremental Clone Detection”. In: *2009 13th European Conference on Software Maintenance and Reengineering*. 2009, pp. 219–228. DOI: 10.1109/CSMR.2009.20.
- [13] D. S. Hirschberg. “A Linear Space Algorithm for Computing Maximal Common Subsequences”. In: *Commun. ACM* 18.6 (June 1975), pp. 341–343. ISSN: 0001-0782. DOI: 10.1145/360825.360861. URL: <https://doi.org/10.1145/360825.360861>.
- [14] Benjamin Hummel et al. “Index-based code clone detection: incremental, distributed, scalable”. In: *2010 IEEE International Conference on Software Maintenance*. 2010, pp. 1–9. DOI: 10.1109/ICSM.2010.5609665.
- [15] Katsuro Inoue. “Introduction to Code Clone Analysis”. In: *Code Clone Analysis: Research, Tools, and Practices*. Ed. by Katsuro Inoue and Chanchal K. Roy. Singapore: Springer Singapore, 2021, pp. 3–27. ISBN: 978-981-16-1927-4. DOI: 10.1007/978-981-16-1927-4_1. URL: https://doi.org/10.1007/978-981-16-1927-4_1.
- [16] G. Jacobson. “Space-efficient static trees and graphs”. In: *30th Annual Symposium on Foundations of Computer Science*. 1989, pp. 549–554. DOI: 10.1109/SFCS.1989.63533.
- [17] Lingxiao Jiang et al. “DECKARD: Scalable and Accurate Tree-Based Detection of Code Clones”. In: *29th International Conference on Software Engineering (ICSE’07)*. 2007, pp. 96–105. DOI: 10.1109/ICSE.2007.30.
- [18] Stephen H. Kan. *Metrics and Models in Software Quality Engineering*. 2nd. USA: Addison-Wesley Longman Publishing Co., Inc., 2002. ISBN: 0201729156.
- [19] Juha Kärkkäinen. “Suffix Array Construction”. In: *Encyclopedia of Algorithms*. Ed. by Ming-Yang Kao. New York, NY: Springer New York, 2016, pp. 2141–2144. ISBN: 978-1-4939-2864-4. DOI: 10.1007/978-1-4939-2864-4_412. URL: https://doi.org/10.1007/978-1-4939-2864-4_412.
- [20] Shinji Kawaguchi et al. “SHINOBI: A Tool for Automatic Code Clone Detection in the IDE”. In: *16th Working Conference on Reverse Engineering, WCRE 2009, 13-16 October 2009, Lille, France*. Ed. by Andy Zaidman, Giuliano Antoniol, and Stéphane Ducasse. IEEE Computer Society, 2009, pp. 313–314. DOI: 10.1109/WCRE.2009.36. URL: <https://doi.org/10.1109/WCRE.2009.36>.
- [21] Microsoft. *Language Server Protocol*. <https://microsoft.github.io/language-server-protocol/>. Accessed: 2023-02-17.
- [22] Tung Thanh Nguyen et al. “Scalable and incremental clone detection for evolving software”. In: *25th IEEE International Conference on Software Maintenance (ICSM 2009), September 20-26, 2009, Edmonton, Alberta, Canada*. IEEE Computer Society, 2009, pp. 491–494. DOI: 10.1109/ICSM.2009.5306283. URL: <https://doi.org/10.1109/ICSM.2009.5306283>.
- [23] Ge Nong, Sen Zhang, and Wai Hong Chan. “Two Efficient Algorithms for Linear Time Suffix Array Construction”. In: *IEEE Trans. Computers* 60.10 (2011), pp. 1471–1484. DOI: 10.1109/TC.2010.188. URL: <https://doi.org/10.1109/TC.2010.188>.
- [24] Chaiyong Ragkhitwetsagul and Jens Krinke. “Siamese: scalable and incremental code clone search via multiple code representations”. In: *Empir. Softw. Eng.* 24.4 (2019), pp. 2236–2284. DOI: 10.1007/s10664-019-09697-7. URL: <https://doi.org/10.1007/s10664-019-09697-7>.

- [25] M. Rieger, S. Ducasse, and M. Lanza. “Insights into system-wide code duplication”. In: *11th Working Conference on Reverse Engineering*. 2004, pp. 100–109. DOI: 10.1109/WCRE.2004.25.
- [26] Chanchal Kumar Roy, James R. Cordy, and Rainer Koschke. “Comparison and evaluation of code clone detection techniques and tools: A qualitative approach”. In: *Sci. Comput. Program.* 74.7 (2009), pp. 470–495. DOI: 10.1016/j.scico.2009.02.007. URL: <https://doi.org/10.1016/j.scico.2009.02.007>.
- [27] M. Salson et al. “Dynamic extended suffix arrays”. In: *Journal of Discrete Algorithms* 8.2 (2010). Selected papers from the 3rd Algorithms and Complexity in Durham Workshop ACiD 2007, pp. 241–257. ISSN: 1570-8667. DOI: <https://doi.org/10.1016/j.jda.2009.02.007>. URL: <https://www.sciencedirect.com/science/article/pii/S1570866709000343>.
- [28] Jeffrey Svajlenko and Chanchal K. Roy. “BigCloneBench”. In: *Code Clone Analysis*. Ed. by Katsuro Inoue and Chanchal K. Roy. Springer Singapore, 2021, pp. 93–105. DOI: 10.1007/978-981-16-1927-4_7. URL: https://doi.org/10.1007/978-981-16-1927-4_7.
- [29] Md Sharif Uddin, Chanchal K. Roy, and Kevin A. Schneider. “Towards Convenient Management of Software Clone Codes in Practice: An Integrated Approach”. In: *Proceedings of the 25th Annual International Conference on Computer Science and Software Engineering*. CASCON ’15. Markham, Canada: IBM Corp., 2015, pp. 211–220.
- [30] Esko Ukkonen. “On approximate string matching”. In: *Foundations of Computation Theory*. Ed. by Marek Karpinski. Berlin, Heidelberg: Springer Berlin Heidelberg, 1983, pp. 487–495. ISBN: 978-3-540-38682-7.
- [31] Esko Ukkonen. “On-line construction of suffix trees”. In: *Algorithmica* 14.3 (1995), pp. 249–260.
- [32] Robert A. Wagner and Michael J. Fischer. “The String-to-String Correction Problem”. In: *J. ACM* 21.1 (Jan. 1974), pp. 168–173. ISSN: 0004-5411. DOI: 10.1145/321796.321811. URL: <https://doi.org/10.1145/321796.321811>.
- [33] Tim A. Wagner. “Practical Algorithms for Incremental Software Development Environments”. PhD thesis. EECS Department, University of California, Berkeley, Mar. 1998. URL: <http://www2.eecs.berkeley.edu/Pubs/TechRpts/1998/5885.html>.
- [34] Zhipeng Xue et al. “SEED: Semantic Graph based Deep detection for type-4 clone”. In: *CoRR* abs/2109.12079 (2021). arXiv: 2109.12079. URL: <https://arxiv.org/abs/2109.12079>.
- [35] Minhaz F. Zibran and Chanchal K. Roy. “IDE-Based Real-Time Focused Search for near-Miss Clones”. In: *Proceedings of the 27th Annual ACM Symposium on Applied Computing*. SAC ’12. Trento, Italy: Association for Computing Machinery, 2012, pp. 1235–1242. ISBN: 9781450308571. DOI: 10.1145/2245276.2231970. URL: <https://doi.org/10.1145/2245276.2231970>.